

加载MNIST数据

```
In [1]: import keras
(X0, Y0), (X1, Y1) = keras.datasets.mnist.load_data()
```

Using TensorFlow backend.

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>


11493376/11490434 [=====] - 1824s 159us/step

数据展示

```
In [2]: from matplotlib import pyplot as plt
plt.figure()
fig, ax = plt.subplots(2, 5)
ax = ax.flatten()

for i in range(10):
    Im = X0[Y0 == i][0]      # Im = Y0对应的数字 (eg 1) 对应的图片 (1的) 里的第一张
    ax[i].imshow(Im)
plt.show()
```

<Figure size 432x288 with 0 Axes>



产生One-Hot型因变量

```
In [3]: from keras.utils import to_categorical
YY0=to_categorical(Y0)
YY1=to_categorical(Y1)
YY1
#TensorFlow只能read One-Hot类型的Y
```

```
Out[3]: array([[0., 0., 0., ..., 1., 0., 0.],
               [0., 0., 1., ..., 0., 0., 0.],
               [0., 1., 0., ..., 0., 0., 0.],
               ...,
               [0., 0., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.]], dtype=float32)
```

逻辑回归模型

```
In [4]: from keras.layers import Activation, Dense, Flatten, Input
from keras import Model

input_shape=(28,28) #一张图片的size
input_layer=Input(input_shape)
x=input_layer
x=Flatten()(x) #使图片变成1*784的矩阵
x=Dense(10)(x) #全链接层, 最后output要是10以对应10个digits
x=Activation('softmax')(x) #把output变成probability
output_layer=x
model=Model(input_layer,output_layer)
```

In [5]: `model.summary()` #查看模型参数

Model: "model_1"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 28, 28)	0
flatten_1 (Flatten)	(None, 784)	0
dense_1 (Dense)	(None, 10)	7850
activation_1 (Activation)	(None, 10)	0

Total params: 7,850
 Trainable params: 7,850
 Non-trainable params: 0

Dense_1 有 $784 \times 10 (= 7840) + 10 = 7850$ 个参数 (weights)。因为在全连接层中每一个从 `flatten_1` 来的都要连接到10个nodes, 所以784 10; 剩下的10个weights来自截距项

In [6]: `from keras.optimizers import Adam`
`model.compile(optimizer = Adam(0.01),`
`loss = 'categorical_crossentropy', #LR比较常用的的loss`
`function`
`metrics = ['accuracy'])`

In [7]: `model.fit(X0,YY0,`
`validation_data=(X1,YY1),`
`batch_size=1000,`
`epochs=20)`

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 1s 15us/step - loss: 27.4137 - accuracy: 0.8112 - val_loss: 9.9834 - val_accuracy: 0.9027

Epoch 2/20

60000/60000 [=====] - 0s 6us/step - loss: 7.7616 - accuracy: 0.9006 - val_loss: 7.0122 - val_accuracy: 0.8939

Epoch 3/20

60000/60000 [=====] - 0s 6us/step - loss: 7.7616 - accuracy: 0.9006 - val_loss: 7.0122 - val_accuracy: 0.8939

```
ss: 6.2137 - accuracy: 0.8925 - val_loss: 5.8425 - val_accuracy
: 0.8942
Epoch 4/20
60000/60000 [=====] - 0s 6us/step - lo
ss: 5.3354 - accuracy: 0.8907 - val_loss: 5.6379 - val_accuracy
: 0.8954
Epoch 5/20
60000/60000 [=====] - 0s 6us/step - lo
ss: 5.1785 - accuracy: 0.8890 - val_loss: 5.6164 - val_accuracy
: 0.8858
Epoch 6/20
60000/60000 [=====] - 0s 6us/step - lo
ss: 5.0149 - accuracy: 0.8892 - val_loss: 6.9368 - val_accuracy
: 0.8712
Epoch 7/20
60000/60000 [=====] - 0s 6us/step - lo
ss: 5.4741 - accuracy: 0.8859 - val_loss: 5.2232 - val_accuracy
: 0.9061
Epoch 8/20
60000/60000 [=====] - 0s 5us/step - lo
ss: 5.5999 - accuracy: 0.8910 - val_loss: 5.7054 - val_accuracy
: 0.9019
Epoch 9/20
60000/60000 [=====] - 0s 6us/step - lo
ss: 5.2817 - accuracy: 0.8910 - val_loss: 6.8126 - val_accuracy
: 0.8707
Epoch 10/20
60000/60000 [=====] - 0s 6us/step - lo
ss: 5.4920 - accuracy: 0.8888 - val_loss: 6.5173 - val_accuracy
: 0.8966
Epoch 11/20
60000/60000 [=====] - 0s 6us/step - lo
ss: 5.1358 - accuracy: 0.8943 - val_loss: 6.5194 - val_accuracy
: 0.8836
Epoch 12/20
60000/60000 [=====] - 0s 6us/step - lo
ss: 5.4373 - accuracy: 0.8903 - val_loss: 6.5411 - val_accuracy
: 0.8889
Epoch 13/20
60000/60000 [=====] - 0s 6us/step - lo
ss: 4.8872 - accuracy: 0.8974 - val_loss: 6.4614 - val_accuracy
: 0.8851
Epoch 14/20
60000/60000 [=====] - 0s 5us/step - lo
ss: 4.8789 - accuracy: 0.8943 - val_loss: 7.3454 - val_accuracy
: 0.8722
Epoch 15/20
60000/60000 [=====] - 0s 6us/step - lo
ss: 5.8168 - accuracy: 0.8894 - val_loss: 6.6594 - val_accuracy
: 0.8972
Epoch 16/20
60000/60000 [=====] - 0s 5us/step - lo
ss: 5.4283 - accuracy: 0.8941 - val_loss: 6.8719 - val_accuracy
: 0.8802
Epoch 17/20
```


```
60000/60000 [=====] - 0s 5us/step - loss: 5.3954 - accuracy: 0.8921 - val_loss: 7.3211 - val_accuracy: 0.8730
Epoch 18/20
60000/60000 [=====] - 0s 6us/step - loss: 5.5083 - accuracy: 0.8942 - val_loss: 6.6847 - val_accuracy: 0.8836
Epoch 19/20
60000/60000 [=====] - 0s 5us/step - loss: 6.4071 - accuracy: 0.8862 - val_loss: 8.1961 - val_accuracy: 0.8879
Epoch 20/20
60000/60000 [=====] - 0s 6us/step - loss: 5.5712 - accuracy: 0.8956 - val_loss: 9.4722 - val_accuracy: 0.8706
```

```
Out[7]: <keras.callbacks.callbacks.History at 0x7f783c22bd50>
```

得到accuracy在0.9左右

参数估计结果可视化

```
In [8]: fig,ax = plt.subplots(2,5)
ax=ax.flatten()
weights = model.layers[2].get_weights()[0]
for i in range(10):
    Im=weights[:,i].reshape((28,28))
    ax[i].imshow(Im,cmap='seismic')
    ax[i].set_title("{}".format(i))
    ax[i].set_xticks([])
    ax[i].set_yticks([])
plt.show()
```



好像并没能看出digits间特别大的差异，但是有一些轮廓可以大概看清。

改进模型

```
In [9]: input_shape=(28,28) #一张图片的size
input_layer=Input(input_shape)
x=input_layer
x=Flatten()(x) #使图片变成1*784的矩阵
x=Dense(70)(x) #全连接层, 目的是为了先根据一些digits间的common feature进行简单分类
x=Dense(10)(x) #全连接层, 最后output要是10以对应10个digits
x=Activation('softmax')(x) #把output变成probability
output_layer=x
model1=Model(input_layer,output_layer)
```

```
In [10]: model1.compile(optimizer = Adam(0.01),
                        loss = 'categorical_crossentropy', #LR比较常用的的loss function
                        metrics = ['accuracy'])
```

```
In [11]: model1.fit(X0,YY0,
                    validation_data=(X1,YY1),
                    batch_size=1000,
                    epochs=20)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 1s 9us/step - loss: 86.3689 - accuracy: 0.7965 - val_loss: 19.4118 - val_accuracy: 0.8983

Epoch 2/20

60000/60000 [=====] - 0s 6us/step - loss: 11.9134 - accuracy: 0.9011 - val_loss: 8.4536 - val_accuracy: 0.8981

Epoch 3/20

60000/60000 [=====] - 0s 6us/step - loss: 6.1554 - accuracy: 0.8946 - val_loss: 6.2628 - val_accuracy: 0.8694

Epoch 4/20

60000/60000 [=====] - 0s 6us/step - loss: 4.1176 - accuracy: 0.8895 - val_loss: 4.2979 - val_accuracy: 0.8793

Epoch 5/20

60000/60000 [=====] - 0s 6us/step - loss: 3.2036 - accuracy: 0.8908 - val_loss: 3.2912 - val_accuracy: 0.8733

Epoch 6/20

60000/60000 [=====] - 0s 6us/step - loss: 2.3869 - accuracy: 0.8894 - val_loss: 2.2883 - val_accuracy: 0.8946

Epoch 7/20

60000/60000 [=====] - 0s 6us/step - loss: 1.7733 - accuracy: 0.8926 - val_loss: 2.0455 - val_accuracy

```
: 0.8887
Epoch 8/20
60000/60000 [=====] - 0s 6us/step - loss: 1.4598 - accuracy: 0.8906 - val_loss: 1.4061 - val_accuracy: 0.8962
Epoch 9/20
60000/60000 [=====] - 0s 6us/step - loss: 1.1061 - accuracy: 0.8931 - val_loss: 1.2074 - val_accuracy: 0.8827
Epoch 10/20
60000/60000 [=====] - 0s 6us/step - loss: 1.0735 - accuracy: 0.8858 - val_loss: 1.2105 - val_accuracy: 0.8816
Epoch 11/20
60000/60000 [=====] - 0s 7us/step - loss: 0.8356 - accuracy: 0.8917 - val_loss: 0.8572 - val_accuracy: 0.8839
Epoch 12/20
60000/60000 [=====] - 0s 6us/step - loss: 0.6454 - accuracy: 0.8941 - val_loss: 0.9397 - val_accuracy: 0.8774
Epoch 13/20
60000/60000 [=====] - 0s 6us/step - loss: 0.5621 - accuracy: 0.8995 - val_loss: 0.6272 - val_accuracy: 0.8910
Epoch 14/20
60000/60000 [=====] - 0s 6us/step - loss: 0.4820 - accuracy: 0.9003 - val_loss: 0.5183 - val_accuracy: 0.9039
Epoch 15/20
60000/60000 [=====] - 0s 6us/step - loss: 0.4569 - accuracy: 0.9004 - val_loss: 0.5136 - val_accuracy: 0.9008
Epoch 16/20
60000/60000 [=====] - 0s 6us/step - loss: 0.4213 - accuracy: 0.9036 - val_loss: 0.5154 - val_accuracy: 0.8923
Epoch 17/20
60000/60000 [=====] - 0s 6us/step - loss: 0.4246 - accuracy: 0.9027 - val_loss: 0.4950 - val_accuracy: 0.8998
Epoch 18/20
60000/60000 [=====] - 0s 6us/step - loss: 0.4054 - accuracy: 0.9043 - val_loss: 0.4534 - val_accuracy: 0.9036
Epoch 19/20
60000/60000 [=====] - 0s 6us/step - loss: 0.3724 - accuracy: 0.9072 - val_loss: 0.4359 - val_accuracy: 0.8981
Epoch 20/20
60000/60000 [=====] - 0s 6us/step - loss: 0.3763 - accuracy: 0.9051 - val_loss: 0.4625 - val_accuracy: 0.8966
```

Out[11]: <keras.callbacks.callbacks.History at 0x7f783c3723d0>

可以看到accuracy依旧在9.0左右，但是从0.89 上升到了0.90

In [16]: `model1.summary()`

Model: "model_2"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	(None, 28, 28)	0
flatten_2 (Flatten)	(None, 784)	0
dense_2 (Dense)	(None, 70)	54950
dense_3 (Dense)	(None, 10)	710
activation_2 (Activation)	(None, 10)	0

Total params: 55,660
 Trainable params: 55,660
 Non-trainable params: 0

In [20]: `#dense_2 Param: (28 * 28 =) 784 * 70 (= 54880) + 70 = 54950`
`#dense_3 Param: 70*10 + 10 = 710`

In [13]: `model1.fit(X0,YY0,`
`validation_data=(X1,YY1),`
`batch_size=5000,`
`epochs=20)`
`#更改了batch_size, 因为55000个data 一个batch只取1000有点少`

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 0s 3us/step - loss: 0.2635 - accuracy: 0.9260 - val_loss: 0.3451 - val_accuracy: 0.9154

Epoch 2/20

60000/60000 [=====] - 0s 3us/step - loss: 0.2419 - accuracy: 0.9326 - val_loss: 0.3220 - val_accuracy: 0.9204

Epoch 3/20

60000/60000 [=====] - 0s 3us/step - loss: 0.2347 - accuracy: 0.9353 - val_loss: 0.3164 - val_accuracy


```
: 0.9210
Epoch 4/20
60000/60000 [=====] - 0s 3us/step - loss: 0.2328 - accuracy: 0.9352 - val_loss: 0.3240 - val_accuracy: 0.9196
Epoch 5/20
60000/60000 [=====] - 0s 3us/step - loss: 0.2348 - accuracy: 0.9341 - val_loss: 0.3219 - val_accuracy: 0.9214
Epoch 6/20
60000/60000 [=====] - 0s 3us/step - loss: 0.2349 - accuracy: 0.9344 - val_loss: 0.3212 - val_accuracy: 0.9202
Epoch 7/20
60000/60000 [=====] - 0s 3us/step - loss: 0.2324 - accuracy: 0.9347 - val_loss: 0.3236 - val_accuracy: 0.9204
Epoch 8/20
60000/60000 [=====] - 0s 3us/step - loss: 0.2331 - accuracy: 0.9349 - val_loss: 0.3220 - val_accuracy: 0.9206
Epoch 9/20
60000/60000 [=====] - 0s 3us/step - loss: 0.2322 - accuracy: 0.9355 - val_loss: 0.3220 - val_accuracy: 0.9204
Epoch 10/20
60000/60000 [=====] - 0s 3us/step - loss: 0.2321 - accuracy: 0.9355 - val_loss: 0.3199 - val_accuracy: 0.9216
Epoch 11/20
60000/60000 [=====] - 0s 3us/step - loss: 0.2314 - accuracy: 0.9355 - val_loss: 0.3195 - val_accuracy: 0.9218
Epoch 12/20
60000/60000 [=====] - 0s 3us/step - loss: 0.2333 - accuracy: 0.9351 - val_loss: 0.3203 - val_accuracy: 0.9214
Epoch 13/20
60000/60000 [=====] - 0s 3us/step - loss: 0.2311 - accuracy: 0.9357 - val_loss: 0.3203 - val_accuracy: 0.9187
Epoch 14/20
60000/60000 [=====] - 0s 3us/step - loss: 0.2315 - accuracy: 0.9351 - val_loss: 0.3151 - val_accuracy: 0.9217
Epoch 15/20
60000/60000 [=====] - 0s 3us/step - loss: 0.2306 - accuracy: 0.9352 - val_loss: 0.3201 - val_accuracy: 0.9220
Epoch 16/20
60000/60000 [=====] - 0s 3us/step - loss: 0.2326 - accuracy: 0.9348 - val_loss: 0.3277 - val_accuracy: 0.9192
Epoch 17/20
60000/60000 [=====] - 0s 3us/step - loss:
```

```
ss: 0.2344 - accuracy: 0.9338 - val_loss: 0.3196 - val_accuracy
: 0.9214
Epoch 18/20
60000/60000 [=====] - 0s 3us/step - lo
ss: 0.2336 - accuracy: 0.9346 - val_loss: 0.3256 - val_accuracy
: 0.9183
Epoch 19/20
60000/60000 [=====] - 0s 3us/step - lo
ss: 0.2310 - accuracy: 0.9357 - val_loss: 0.3280 - val_accuracy
: 0.9188
Epoch 20/20
60000/60000 [=====] - 0s 3us/step - lo
ss: 0.2342 - accuracy: 0.9345 - val_loss: 0.3261 - val_accuracy
: 0.9185
```

```
Out[13]: <keras.callbacks.callbacks.History at 0x7f78300aaf50>
```

accruacy 上升到了0.93

更多基于图片的分类问题

应用一 动物分类

例子： 狗狗品种分类， Dog or muffin 等
(以 dog or muffin 举例)

X: 一张要么是狗要么是 muffin 的图片

Y: 狗或 muffin

应用二 植物分类

例子： 识别不同植物种类的 APP

X: 对一个植物拍照 -> 提取有关植物的信息 (叶片形状, 茎长宽, 颜色等)

Y: 与 X 信息相匹配的植物名称

应用三 物品识别

例子： 垃圾分类

X: 一张垃圾的图片

Y: 可回收物/其他垃圾/有害垃圾等垃圾的分类